

1 Heat Equation - 1D

1.1 Round Robin Decomposition:

Question: Write the algorithm that computes heat equation using a Round Robin decomposition.

CalculateHeat1D(N, Heat_K, p, P){

```
if(N%P != 0){
if(p == P-1) {
end = N;
}
}

for(i = p; i < N; i+=p) { //for loop that starts working depending on its process id and is incremented
                           by total number of processors to implement round robin process
Heat_K-1[i] = GenerateHeat_K-1(); //Function to generate process specific Heat_K-1 element.

if(i == 0) // condition to check if its the start of the array
{
Isend Heat_K-1[i] to p+1
Irecv Heat_K-1[i+1] from p+1
Heat_K[i] = (2*Heat_K-1[i] +Heat_K-1[i+1])/3;
}
else if ( i== N-1) // condition to check if its the end of the array
{
Isend Heat_K-1[i] to p-1
Irecv Heat_K-1[i-1] from p-1
Heat_K[i] = (2*Heat_K-1[i] + Heat_K-1[i-1])/3;
}
else
{

Isend Heat_K-1[i] to p-1
Isend Heat_K-1[i] to p+1
Irecv Heat_K-1[i+1] from p+1
Irecv Heat_K-1[i-1] from p-1
Heat_K[i] = (Heat_K-1[i-1] + Heat_K-1[i] +Heat_K-1[i+1])/3;
}
}
}
```

Question: How much communication happen per iteration of the heat equation for a Round Robin decomposition?

- $O(N)$

1.2 Block Decomposition

Question: Write the algorithm that computes heat equation using a Block decomposition.

-

```
CalculateHeat1D( N, Heat_K, p, P){

block_size = N/P;
start = p * block_size
end = (p+1) * block_size

if(N%P != 0){ // if N/P is not divisible entirely by P processors then the last processor is assigned the
                extra elements
if(p == P-1) {
end = N;
}
}

arrsize = end - start
Heat_K-1[arrsize]; //creating array K-1 of required size depending on the processid.

for(i = start ; i< end; i++) { //for loop runs from its corresponding start to end elements

Heat_K-1[i] = GenerateHeat_K-1(); //Function to generate process specific Heat_K-1 element.

if(i == 0) // condition to check if its the start of the array
{
Isend Heat_K-1[i] to p+1
Recv Heat_K-1[i+1] from p+1
Heat_K[i] = (2*Heat_K-1[i] +Heat_K-1[i+1])/3;
}
else if ( i== N-1) // condition to check if its the end of the array
{
Isend Heat_K-1[i] to p-1
Irecv Heat_K-1[i-1] from p-1
Heat_K[i] = (2*Heat_K-1[i] + Heay_K-1[i-1])/3;
}
else if ( i== p*(block_size)) // condition to check if the element is the start element of any process
                                other than zero inorder to get the value Heat_K-1[i-1] from previous value and
                                give Heat_K-1[i] to p-1
{
Isend Heat_K-1[i] to p-1
Irecv Heat_K-1[i-1] from p-1
Heat_K[i] = (Heat_K-1[i-1] + Heat_K-1[i] +Heay_K-1[i-1])/3;
}
else if ( i== end)
```

```

{
  Isend Heat_K-1[i] to p+1
  Irecv Heat_K-1[i+1] from p+1
  Heat_K[i] = (Heat_K-1[i-1] + Heat_K-1[i] + Heat_K-1[i+1])/3;
}
else
{
  Heat_K[i] = (Heat_K-1[i-1] + Heat_K-1[i] + Heat_K-1[i+1])/3;
}
}

```

Question: How much communication happen per iteration of the heat equation when using a Block decomposition?

- $O(N)$

1.3 Reflection

Question: What data partitioning would you use?

- I would use Block Partitioning over Roundrobin

2.1 1D partitioning: Horizontal stripes:

Question: Write the algorithm that performs $y = Ax$; $x = y$; 10 times in a loop if the data is partitioned horizontally.

MatVec(N, X[], p, P)

```

{
  block_size = N/P;
  start = p* block_size ;
  end = (p+1) * block_size ;

```

```

if(N%P != 0)

```

```

{
  if(p == P-1)
  {
    end = N;
  }
}

```

```

rows = end - start;

```

```

int[][] A = new int[rows][N];

```

```

int[] Y= new int[rows];

```

```

for(int i= start ; i< end; i++)
{
    for(int j= 0; j< N; j++)
    {
        A[i][j] = Generate();
    }
}

```

```

for(int iter = 0; iter < 10; iter++)
{

```

```

for(int i= start ; i< end; i++)
{
    for(int j= 0; j< N; j++)
    {
        Y[i] += A[i][j] *X[j];
    }
}

```

using for loop starting from i= 0 to P-1

send the data to all the other processes except for the process itself which can be done by introducing an

If condition of

```

    if(i != p) {
        ISend(Y,i);
    }

```

Receive the data similarly from all the nodes and replace the existing X with these values depending upon the datacount and source from where the data was sent.

eg. If the data is from 0 process get the data count sent and using for loop like

for i= source*blocksize to (source* blocksize) + datacount

```

for( t = 0 ; t< P; t++){
if(t != p)
{
    Send Y to t;
}
}

```

```

for(r = 0 to p-1)
{
    if (r != p)
    {
        recv Y from r;
        source = r
        for (i= source*blocksize; i< (source* blocksize ) + datacount; i++)
        {
            memcpy received Y to X
        }
    }
}

```

```
}  
}
```

Question: How much memory does each node need if the data is partitioned horizontally?

-
 N^2

Question: How much communication does the algorithm do per iteration if the data is partitioned horizontally

-
 $O(N)$

2.2 1D partitioning: vertical stripes

Question: Write the algorithm that performs $y = Ax; x = y$; 10 times in a loop if the data is partitioned vertically.

-
MatVec(N, Y[], p, P)
{

 block_size = N/P;

 start = p* block_size ;
 end = (p+1) * block_size ;

 if(N%P != 0)
 {
 if(p == P-1)
 {
 end = N;
 }
 }

 size = end-start;

 int[][] A = new [N][size];
 int[] X= new int[size];

 for(int i= 0; i< N; i++)
 {
 for(int j= start ; j< end; j++)
 {
 A[i][j] = Generate();
 }
 }

 for(int j= start ; j< end; j++)

```
{
  X[j] = Generate();
}
```

```
for(int iter = 0; iter < 10; iter++)
{
  for(int i= 0; i< N; i++)
  {
    for(int j= start ; j< end; j++)
    {
      Y[i] += A[i][j] *X[j];
    }
  }
}
```

using for loop starting from i= 0 to P-1

send the data to all the other processes except for the process itself which can be done by introducing an

If condition of

```
if(i != p) {
  Send(Y, I);
}
```

Receive the data similarly from all the nodes, reduce the values received from all the nodes and replace the existing X with the reduced value.

```
for( t = 0 ; t< P; t++)
{
  if(t != p)
  {
    Send Y to t;
  }
}
```

```
for( t = 0 ; t< P; t++)
{
  if(t != p)
  {
    Recv Y from t;
  }
  for(i=0 ; i< N; i++)
  {
    ReduceAll(Y[i] into X[i]);
  }
}
}
```

Question: How much memory does each node need if the data is partitioned vertically?

-
 N^2

Question: How much communication does the algorithm do per iteration if the data is partitioned vertically?

-
 $O(N)$

2.3 2D partitioning: blocks

Question: Write the algorithm that performs $y = Ax; x = y$; 10 times in a loop if the data is partitioned in blocks.

MatVec(N, Y[], p, P)

```
{  
  
    int square = sqrt(P);  
    int block_size = N/square;  
    int div = p / square;  
    int mod = fmod(p, square);  
    int rstart = (p - (p - div)) * block_size;  
    int rend = rstart + (block_size-1);  
    int cstart = (p - (p - mod)) * block_size;  
    int cend = cstart + (block_size-1);
```

```
    int[][] A = new [block_size][block_size];  
    int[] X= new int[block_size];  
    int[] Y= new int[block_size];
```

```
    for(int i= rstart; i< rend; i++)  
    {  
        for(int j= cstart ; j< cend; j++)  
        {  
            A[i][j] = Generate();  
        }  
    }
```

```
    for(int j= colstart ; j< colend; j++)  
    {  
        X[j] = Generate();  
    }
```

```
    for(int iter = 0; iter < 10; iter++)  
    {  
        for(int i= rstart; i< rend; i++)  
        {
```

```

for(int j= cstart ; j< cend; j++)
{
    Y[i] += A[i][j] *X[j];
}
}

```

```

fake_p = p / square; //to get the exact row in which te data has to be sent
//finding the processor id s to which the data needs to be sent
start_elem = fake_p * square;
end_elem = start_elem + square;

```

```

for(e=start_elem ; e<end_elem; e++)
{
    if(e != p)
    {
        send Yto e; // send to the processes here given by e
    }
}

```

```

for(e=start_elem ; e<end_elem; e++)
{
    if(e != p)
    {
        Recv Yfrom e;
        for(f = 0; f< datacount; f++) //datacount is the number of elements received
        {
            Reduceall(Y[f] to X[f]); Reduce the value in X received in Y
        }
    }
}
}
}

```

Question: How much memory does each node need if the data is partitioned in blocks?

- N^2

Question: How much communication does the algorithm do per iteration if the data is partitioned in blocks?

- $O(N)$

3 Reduction

Question: Fill the following table. For each algorithm and each network structure, answer the following questions. Run a small example if you have difficulty seeing how communication happens; but express all answers for the case with P processors.

Case	How much data on most loaded link	How much data on most loaded node	How long is the longest chain of communication.
Reduce-star on chain	$O(P-1)$	$O(P-1)$	P
Reduce-star on clique	$O(1)$	$O(P-1)$	1
Reduce-chain on chain	$O(1)$	$O(1)$	P
Reduce-chain on clique	$O(1)$	$O(1)$	P
Reduce-tree on chain	$O(P/2)$	$O(\log(P))$	$P/2$
Reduce-tree on clique	$O(1)$	$O(\log(P))$	1

Question: What do you think is the best algorithm for each network structure? (One of the given algorithm or a different one.)

- Among the above algorithms the best fitting algorithm is Reduce-Tree Algorithm.